
Search-Guided Evolutionary Optimization for Graph Pathfinding

A comparative study of classical and hybrid search methods

Author: Naga Sampath Maddineni
Project type: Reproducible course project

Abstract

This project studies route planning on deterministic obstacle grids through a common experimental interface. Six classical methods, including uninformed, cost-based, and heuristic search, are compared with a Genetic Algorithm. A hybrid variant seeds the evolutionary population with an A* route. The benchmark contains 12 generated scenarios spanning three grid sizes, two obstacle densities, and unit-cost or weighted terrain. Across the executed trials, classical methods reached every generated goal. On weighted maps, UCS and A* recorded the lowest mean path cost among the classical methods. Across 120 evolutionary trials, random initialization succeeded in 73.3% of trials, while A*-seeded initialization succeeded in 100%. The findings support a narrow conclusion for this benchmark: search-derived initialization improved evolutionary reliability and final fitness, while classical algorithms retained clearer optimality guarantees.

Keywords

Graph search; A*; Dijkstra; Genetic Algorithm; pathfinding; evolutionary optimization; reproducibility

Project provenance

The repository originally contained two coursework notebooks. HW01 was an eight-row housing-price regression using normalization and gradient descent. HW02 was an MNIST PyTorch MLP. Because neither notebook implemented graph search, the final project preserves them as legacy references and adds a new, self-contained pathfinding study rather than making an unsupported claim about their contents.

1. Introduction and problem statement

Path planning exposes a useful distinction between finding any feasible route and finding a low-cost route efficiently. Breadth-first and depth-first procedures differ in traversal order. Uniform-cost search accounts for edge cost, while heuristic methods use problem structure to reduce search effort. Evolutionary optimization provides a different mechanism: it explores a population of candidate action sequences and can express penalties beyond a single shortest-path objective.

This study asks whether a classical route can provide a useful prior for an evolutionary optimizer without hiding the cost of obtaining that prior. The experimental object is a four-neighbor obstacle grid. Every method receives the same start, goal, obstacles, and terrain costs. The comparison records route quality and the amount of search effort used to obtain it.

2. Research questions and contributions

RQ1. How do uninformed, cost-based, and heuristic methods differ on unit-cost and weighted grids?

RQ2. Does A* seeding change Genetic Algorithm success rate, fitness, convergence, or variability?

RQ3. How do grid size and obstacle density affect the observed trade-offs?

The project contributes an executable notebook, a deterministic benchmark generator, comparable metrics for six classical methods, a random-initialized GA baseline, an A*-seeded hybrid ablation, machine-readable results, and an academic report and presentation generated from the same run.

3. Related work

Dijkstra's shortest-path procedure established a cost-ordered approach for graphs with nonnegative edge costs [1]. Hart, Nilsson, and Raphael formalized heuristic search and the conditions under which A* can retain an optimality property [2]. Holland's work provides the evolutionary basis for selection, crossover, and mutation used by the Genetic Algorithm [3]. The present study uses these established ideas as baselines and evaluates a deliberately narrow hybrid: an A* route becomes an informed member of the initial evolutionary population.

4. System overview




Figure 1. Representative obstacle grid. Dark cells are obstacles; the red line illustrates a feasible route between the green start and red goal. The executed notebook generates the benchmark and records each scenario seed.

The workflow is: generate a reachable grid, run all deterministic baselines, run five trials for each GA initialization, write raw results, then generate figures and summary statistics from those files.

Architecture and formal objective

How the single notebook turns controlled inputs into auditable evidence



Formal definitions

For a decoded route $P = (v_0, v_1, \dots, v_k)$, the measured traversal cost is:

$$C(P) = \sum c(v_i), \quad i = 1 \dots k$$

The evolutionary objective combines route cost with explicit feasibility penalties. For chromosome x :

$$F(x) = C(P_x) + 20d_x + 1.5I_x + 0.75R_x + 0.05|P_x|$$

Here d_x is remaining Manhattan distance for an incomplete route, I_x counts invalid actions, R_x counts revisits, and $|P_x|$ is the decoded path length. Completed routes omit the remaining-distance term. Lower fitness is preferred.

The architecture keeps generation, optimization, measurement, and reporting in one inspectable notebook while preserving a machine-readable audit trail.

5. Benchmark and methodology

The generator creates square four-neighbor grids at 15, 25, and 35 cells per side. It samples obstacle layouts at densities 0.15 and 0.25, reserves the start and goal, and retries with a deterministic seed offset until a route exists. Unit-cost scenarios assign cost 1 to every move. Weighted scenarios assign traversable cells one of the costs 1, 3, or 7. The same grid layout is used to compare the two cost modes.

Design element	Configuration
Scenarios	12 maps: 3 sizes x 2 densities x 2 cost modes
Map seeds	7000 through 7011
GA seeds	11, 22, 33, 44, 55
Population / generations	30 / 50
Elite / crossover / mutation	2 / 0.80 / 0.05
Action budget	6 x grid size

6. Metrics and fairness controls

For classical search, the notebook records success, path length, path cost, nodes expanded, nodes generated, maximum frontier size, and runtime. UCS and A* provide the cost reference on weighted maps because the Manhattan heuristic is admissible when multiplied by the minimum terrain cost. BFS is interpreted as length-optimal only in unit-cost settings.

The GA uses a fixed-length action chromosome. Invalid moves and revisits receive penalties; incomplete chromosomes receive a remaining-distance penalty. Both GA variants use the same population size, generation count, selection, crossover, mutation rate, and five seeds. The hybrid records the A* seed-generation time separately so the initialization trade-off remains visible.

7. Algorithm design

Family	Methods	Primary comparison
Uninformed	BFS, DFS, IDS	Feasibility, length, expansion effort
Cost / heuristic	UCS, Greedy, A*	Weighted cost and search effort
Evolutionary	Random GA	Stochastic route optimization
Hybrid	A*-seeded GA	Effect of informed initialization

8. Classical search results




Figure 2. Mean path cost across the 12-map classical comparison. Weighted terrain separates methods that only minimize steps from methods that account for traversal cost.

Method	Unit success	Unit cost	Weighted success	Weighted cost	Unit expansions
BFS	100%	48.33	100%	184.67	552.5
DFS	100%	73.67	100%	229.67	96.2
IDS	100%	48.33	100%	184.67	114936.0
UCS	100%	48.33	100%	122.00	552.5
Greedy	100%	53.33	100%	153.33	68.2
A*	100%	48.33	100%	122.00	382.7

All classical methods found a route on every generated map. On unit-cost maps, BFS, IDS, UCS, and A* shared the lowest mean path cost of 48.33. On weighted maps, UCS and A* shared the lowest mean cost of 122.00. DFS and Greedy expanded fewer nodes on average in this benchmark, but their routes were more expensive. IDS found the same unit-cost route length as BFS while expanding substantially more nodes because it repeated bounded-depth searches.

9. Evolutionary and hybrid results





Figure 3. Mean best-fitness trajectory across five seeds on the 25x25, 15% unit-cost scenario. Lower fitness is better.

Initialization	All success	Mean final fitness	Std. dev.	Mean successful cost
Random GA	73.3%	565.746	593.212	174.14
A*-seeded GA	100.0%	87.658	50.080	85.17

Across all 120 evolutionary trials, random initialization succeeded in 73.3% of trials and A*-seeded initialization succeeded in 100%. Mean final fitness decreased from 565.746 to 87.658. The seeded variant also matched the classical reference cost on successful unit and weighted trials in the aggregate table. These results support the hybrid hypothesis for this benchmark, while the comparison remains conditional on the chosen grid generator, action encoding, and A* prior.

Scalability and density sensitivity

A* node expansions by grid size and obstacle density, computed from algorithm_comparison.csv



Mean nodes expanded; lower is less search effort. Each bar averages the two cost modes for the selected size and density.

Grid	15% density	25% density
15x15	157.5	155.5
25x25	472.0	352.5
35x35	932.0	687.5

Interpretation

The figure makes the scalability question explicit: larger grids generally increase frontier work, while obstacle density changes the reachable geometry and can either increase or reduce expansions for a particular seed. This is a descriptive benchmark result, not a complexity proof or a universal ranking.

All values are recomputed from the committed CSV rows; runtime remains environment-dependent and is not used as a cross-machine claim.

10. Discussion

The classical results reflect the objectives built into each algorithm. BFS and IDS are aligned with step count on unit-cost grids, but they do not account for weighted terrain. UCS and A* optimize the weighted objective, and A* reduces the amount of expansion relative to UCS on unit-cost maps in the aggregate results. Greedy Best-First Search uses the same geometric signal without the accumulated-cost guarantee, which explains its lower expansion count alongside higher path cost.

The GA comparison exposes a different trade-off. Random chromosomes must discover feasible action sequences through mutation and selection, so some trials remain incomplete after the fixed budget. A* seeding places a feasible, low-cost route in the initial population and preserves it through elitism. This directly improves reliability in the measured benchmark, but it also means the hybrid inherits information from a classical solver. The result should be interpreted as evidence for search-guided initialization, not as evidence that an uninformed GA dominates exact search.

11. Limitations and threats to validity

- The benchmark is synthetic and uses small-to-medium grids, so it does not represent all graph topologies or real routing environments.
- The GA uses a simple fixed-length action chromosome. A different encoding, repair policy, population size, or fitness weighting could change the result.
- Runtime measurements are environment-dependent and should be compared within the same execution environment.
- The seeded method receives an A* route by design. Its benefit is therefore a hybridization result, not an independent comparison against a purely evolutionary method.
- The experiment reports aggregate means and standard deviations but does not claim statistical generalization beyond the generated scenarios.

12. Future work

Future extensions could evaluate irregular graphs, multi-objective route quality, larger repeated benchmark suites, alternative evolutionary encodings, and a cost-aware repair operator. A stronger study would also pre-register scenario seeds and compare confidence intervals across independently generated map families.

13. Reproducibility and repository artifacts

The executed notebook is the single project implementation. It generates the graph instances, runs every algorithm, writes the result files, and creates the SVG evidence. The repository records the exact scenario configuration in `experiment_manifest.json`. No restricted course dataset is required.

Artifact	Purpose
<code>src/Graph_Search_Evolutionary_Optimization.ipynb</code>	Complete executed implementation and narrative
<code>results/algorithm_comparison.csv</code>	72 deterministic classical-search rows
<code>results/stochastic_trials.csv</code>	120 GA trial rows with seeds and fitness
<code>results/convergence.csv</code>	6,000 generation-level convergence rows
<code>results/figures/</code>	Generated path, comparison, and convergence SVGs
<code>reports/...Report.pdf</code>	Academic project report
<code>reports/...Presentation.pptx</code>	Academic presentation

14. Conclusion

This study converted a repository of disconnected coursework artifacts into a reproducible search-and-optimization project. Classical search methods provided dependable reference solutions, with UCS and A* producing the lowest weighted path cost in the executed benchmark. The A*-seeded Genetic Algorithm improved evolutionary reliability from 73.3% to 100% successful trials and reduced mean final fitness from 565.746 to 87.658. The conclusion is intentionally bounded: informed initialization helped on these deterministic grids, while exact classical methods remained the clearest reference for optimal path cost.

References

- [1] E. W. Dijkstra, "A note on two problems in connexion with graphs," *Numerische Mathematik*, vol. 1, pp. 269–271, 1959. doi: 10.1007/BF01386390.
- [2] P. E. Hart, N. J. Nilsson, and B. Raphael, "A formal basis for the heuristic determination of minimum cost paths," *IEEE Transactions on Systems Science and Cybernetics*, vol. 4, no. 2, pp. 100–107, 1968. doi: 10.1109/TSSC.1968.300136.
- [3] J. H. Holland, *Adaptation in Natural and Artificial Systems*, University of Michigan Press, 1975. doi: 10.7551/MITPRESS/1090.001.0001.